

Tanzanian Water Wells Prediction

Authors: JOSPHAT KIMANI,

1. Business Understanding

Problem Statement

The purpose of this project is to address the challenge faced by Wells of Life, an NGO operating in Tanzania, in effectively identifying water wells in need of repair. With over 57 million people struggling to access clean water, many existing water points require maintenance or have failed entirely. The problem statement revolves around the inefficiency in prioritizing maintenance efforts due to limited resources and the lack of a systematic approach to assess the condition of water wells. To tackle this issue, the project aims to develop a classifier that utilizes various data points such as pump type, installation date, and other relevant variables to predict the condition of water wells accurately. By doing so, Wells of Life can optimize their resource allocation, prioritize critical repairs, and intervene proactively to prevent disruptions in water supply, ultimately improving access to clean water for communities in Tanzania.

Specific Objectives

- Develop insights into trends and patterns distinguishing between non-functional and functional wells.
- Utilize straightforward analysis to pinpoint non-functioning wells and forecast well functionality based on available variables.

Research Questions

- How can historical data accurately predict water well conditions?
- What challenges exist in predicting well conditions, and how can they be addressed?
- How does the classifier's performance compare to existing methods?
- What insights can the classifier provide for decision-making and resource allocation?

Success Metrics

- To ensure that newly constructed wells are of good quality water for the communities.
- To correctly identify functionality of a well and determine its viability.
- Generating a model that will be able to correctly predict the quality status of the wells in Tanzania with an accuracy of 80%.

2. Data Understanding

- Data Source

The provided dataset from Taarifa and the Tanzanian Ministry of Water is instrumental in predicting the functionality of water pumps in Tanzania.

- Describe Data

By analyzing various factors such as pump type, installation date, and management practices, the dataset allows for the classification of pumps into three categories: functional, in need of repair, or non-functional. This intermediate-level practice competition aims to leverage the dataset to enhance maintenance operations and ensure access to clean water for Tanzanian communities. A thorough understanding of the dataset's source, properties, and suitability for addressing the real-world problem is essential for developing accurate predictive models. Overall, the dataset presents an opportunity to apply machine learning techniques to improve water infrastructure management and enhance water accessibility in Tanzania.

Load Libraries

```
# Loading all the neccessary modules
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn import datasets
from sklearn.model_selection import train_test_split, cross_validate,
GridSearchCV
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from xgboost import XGBClassifier
from category_encoders import TargetEncoder
import pickle
import joblib
import warnings
plt.style.use('fivethirtyeight')
sns.set_style('whitegrid')
pd.set_option('display.max_columns', 40)
warnings.filterwarnings('ignore')
```

Load Data

```
# Loading well values as a pandas DataFrame
df_values = pd.read_csv("./Data/well_data_values.csv", index_col=0)

# Loading well labels as a pandas DataFrame
df_labels = pd.read_csv("./Data/well_data_labels.csv", index_col=0)

# Print shape of the two DataFrames
df_values.shape
```

```
df_labels.shape
```

```
# Create a combined DataFrame with both the value and label data  
joined on the 'id' column
```

```
data = df_values.join(df_labels, on='id')
```

```
data.shape
```

```
# Preview DataFrame
```

```
data.head(5)
```

	amount_tsh	date_recorded	funder	gps_height
installer \				
id				
69572	6000.0	2011-03-14	Roman	1390
Roman				
8776	0.0	2013-03-06	Grumeti	1399
GRUMETI				
34310	25.0	2013-02-25	Lottery Club	686
World				
vision				
67743	0.0	2013-01-28	Unicef	263
UNICEF				
19728	0.0	2011-07-13	Action In A	0
Artisan				

	longitude	latitude	wpt_name	num_private
id				
69572	34.938093	-9.856322	none	0
8776	34.698766	-2.147466	Zahanati	0
34310	37.460664	-3.821329	Kwa Mahundi	0
67743	38.486161	-11.155298	Zahanati Ya Nanyumbu	0
19728	31.130847	-1.825359	Shuleni	0

	basin	subvillage	region	region_code
id				
69572	Lake Nyasa	Mnyusi B	Iringa	11
8776	Lake Victoria	Nyamara	Mara	20
34310	Pangani	Majengo	Manyara	21
67743	Ruvuma / Southern Coast	Mahakamani	Mtwara	90
19728	Lake Victoria	Kyanyamisa	Kagera	18

	district_code	lga	ward	population	public_meeting
\					
id					
69572	5	Ludewa	Mundindi	109	True
8776	2	Serengeti	Natta	280	NaN

34310	4	Simanjiro	Ngorika	250	True
67743	63	Nanyumbu	Nanyumbu	58	True
19728	1	Karagwe	Nyakasimbi	0	True
recorded_by scheme_management					
scheme_name \					
id					
69572	GeoData Consultants Ltd	VWC			
Roman					
8776	GeoData Consultants Ltd	Other			
NaN					
34310	GeoData Consultants Ltd	VWC	Nyumba ya mungu pipe		
scheme					
67743	GeoData Consultants Ltd	VWC			
NaN					
19728	GeoData Consultants Ltd	NaN			
NaN					
permit construction_year extraction_type extraction_type_group					
\					
id					
69572	False	1999	gravity	gravity	
8776	True	2010	gravity	gravity	
34310	True	2009	gravity	gravity	
67743	True	1986	submersible	submersible	
19728	True	0	gravity	gravity	
extraction_type_class management management_group					
payment \					
id					
69572	gravity	vwc	user-group	pay	
annually					
8776	gravity	wug	user-group	never	
pay					
34310	gravity	vwc	user-group	pay per	
bucket					
67743	submersible	vwc	user-group	never	
pay					
19728	gravity	other	other	never	
pay					

quantity_group \ id	payment_type	water_quality	quality_group	quantity
---------------------	--------------	---------------	---------------	----------

69572	annually	soft	good	enough
enough				
8776	never pay	soft	good	insufficient
insufficient				
34310	per bucket	soft	good	enough
enough				
67743	never pay	soft	good	dry
dry				
19728	never pay	soft	good	seasonal
seasonal				

id	source	source_type	source_class \
69572	spring	spring	groundwater
8776	rainwater harvesting	rainwater harvesting	surface
34310	dam	dam	surface
67743	machine dbh	borehole	groundwater
19728	rainwater harvesting	rainwater harvesting	surface

status_group \ id	waterpoint_type	waterpoint_type_group
69572	communal standpipe	communal standpipe
functional		
8776	communal standpipe	communal standpipe
functional		
34310	communal standpipe multiple	communal standpipe
functional		
67743	communal standpipe multiple	communal standpipe non
functional		
19728	communal standpipe	communal standpipe
functional		

```
# View info of the dataframe
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 59400 entries, 69572 to 26348
Data columns (total 40 columns):
#   Column                Non-Null Count  Dtype
---  -
0   amount_tsh            59400 non-null  float64
1   date_recorded         59400 non-null  object
2   funder                55763 non-null  object
```

3	gps_height	59400	non-null	int64
4	installer	55745	non-null	object
5	longitude	59400	non-null	float64
6	latitude	59400	non-null	float64
7	wpt_name	59398	non-null	object
8	num_private	59400	non-null	int64
9	basin	59400	non-null	object
10	subvillage	59029	non-null	object
11	region	59400	non-null	object
12	region_code	59400	non-null	int64
13	district_code	59400	non-null	int64
14	lga	59400	non-null	object
15	ward	59400	non-null	object
16	population	59400	non-null	int64
17	public_meeting	56066	non-null	object
18	recorded_by	59400	non-null	object
19	scheme_management	55522	non-null	object
20	scheme_name	30590	non-null	object
21	permit	56344	non-null	object
22	construction_year	59400	non-null	int64
23	extraction_type	59400	non-null	object
24	extraction_type_group	59400	non-null	object
25	extraction_type_class	59400	non-null	object
26	management	59400	non-null	object
27	management_group	59400	non-null	object
28	payment	59400	non-null	object
29	payment_type	59400	non-null	object
30	water_quality	59400	non-null	object
31	quality_group	59400	non-null	object
32	quantity	59400	non-null	object
33	quantity_group	59400	non-null	object
34	source	59400	non-null	object
35	source_type	59400	non-null	object
36	source_class	59400	non-null	object
37	waterpoint_type	59400	non-null	object
38	waterpoint_type_group	59400	non-null	object
39	status_group	59400	non-null	object

dtypes: float64(3), int64(6), object(31)
memory usage: 18.6+ MB

```
# View Numerical Columns
numerical_columns = data.select_dtypes(include='number').columns

# View Categorical Columns
categorical_columns = data.select_dtypes(include='object').columns
```


The dataset contains 40 columns and 59,400 records. Here's a breakdown of the columns:

- **Numerical Columns (9)**
 - `amount_tsh`: Total static head (amount of water available)

- `gps_height`: Altitude of the well
- `longitude`: GPS coordinate (longitude)
- `latitude`: GPS coordinate (latitude)
- `num_private`: Unspecified numerical attribute
- `region_code`: Geographic location code (region)
- `district_code`: Geographic location code (district)
- `population`: Population around the well
- `construction_year`: Year the waterpoint was built
- **Categorical Columns (31)**
 - `date_recorded`: Date the data was recorded
 - `installer`: Organization that installed the well
 - `funder`: Organization that funded the well
 - `wpt_name`: Name of the waterpoint
 - `basin`: Geographic water basin
 - `subvillage`: Specific location within a village
 - `region`: Geographic region
 - `lga`: Local government area
 - `ward`: Administrative ward
 - `public_meeting`: Whether there was a public meeting about the well (True/False)
 - `recorded_by`: Group entering the data
 - `scheme_management`: Entity managing the waterpoint
 - `scheme_name`: Name of the waterpoint scheme
 - `permit`: Whether the waterpoint is permitted (True/False)
 - `extraction_type`: Method of water extraction
 - `extraction_type_group`: Group of water extraction methods
 - `extraction_type_class`: Class of water extraction methods
 - `management`: How the waterpoint is managed
 - `management_group`: Group of management methods
 - `payment`: Payment type for water
 - `payment_type`: Type of payment for water
 - `water_quality`: Quality of the water
 - `quality_group`: Grouped quality of the water
 - `quantity`: Quantity of water available
 - `quantity_group`: Grouped quantity of water available
 - `source`: Water source
 - `source_type`: Type of water source
 - `source_class`: Class of water source
 - `waterpoint_type`: Type of waterpoint
 - `waterpoint_type_group`: Grouped type of waterpoint
 - `status_group`: Condition of the wells (target variable)

```
# Visualize missing data in the DataFrame
```

```
msno.bar(data, figsize=(20, 5), fontsize=10, color='deepskyblue') #  
Changing color to ocean blue  
plt.title("Missing Data Visualization")  
plt.xticks(rotation=90, fontsize=10)  
plt.show()
```



While the dataset as a whole has relatively few missing values, the `scheme_name` column stands out with nearly half of its data missing.

```
# viewing numerical data  
data.describe()
```

	amount_tsh	gps_height	longitude	latitude
num_private \				
count	59400.000000	59400.000000	59400.000000	5.940000e+04
mean	317.650385	668.297239	34.077427	-5.706033e+00
std	2997.574558	693.116350	6.567432	2.946019e+00
min	0.000000	-90.000000	0.000000	-1.164944e+01
25%	0.000000	0.000000	33.090347	-8.540621e+00
50%	0.000000	369.000000	34.908743	-5.021597e+00
75%	20.000000	1319.250000	37.178387	-3.326156e+00
max	350000.000000	2770.000000	40.345193	-2.000000e-08

region_code	district_code	population	construction_year
-------------	---------------	------------	-------------------

count	59400.000000	59400.000000	59400.000000	59400.000000
mean	15.297003	5.629747	179.909983	1300.652475
std	17.587406	9.633649	471.482176	951.620547
min	1.000000	0.000000	0.000000	0.000000
25%	5.000000	2.000000	0.000000	0.000000
50%	12.000000	3.000000	25.000000	1986.000000
75%	17.000000	5.000000	215.000000	2004.000000
max	99.000000	80.000000	30500.000000	2013.000000

The descriptive statistics table provides insights into the numerical attributes of the dataset, including measures of central tendency, variability, and distribution, offering a comprehensive overview of the data's characteristics.

3. Data Preparation

This phase, commonly known as "data wrangling", involves preparing the dataset(s) for modeling by performing tasks such as:

- Data selection
- Data cleaning
- Exploratory Data Analysis (EDA)

3.1 Data Selection

In this section, the aim is to determine the columns from the dataset that will be utilized.

The evaluation will begin by assessing the similarity among the columns, as many of the 40 columns in the dataset appear to be related, before proceeding to evaluate the remaining columns.

3.1.1 `scheme_management` | `management` | `management_group`

Within the context of the project, the columns `scheme_management`, `management`, and `management_group` all relate to the management of the well. An examination of their value counts will be conducted to ascertain the similarity of the data.

```
def display_value_counts(data, columns):
    """
    Display value counts of specified columns in a DataFrame.

    Parameters:
    - data: DataFrame
        The DataFrame containing the data.
    - columns: list
        A list of column names for which value counts are to be
        displayed.
    """
    for column in columns:
        print(f"Value counts for {column}:")
```

```

print(data[column].value_counts())
print("\n")

# Define columns to analyze
columns_to_analyze = ['scheme_management', 'management',
'management_group']

# Call display_value_counts function
display_value_counts(data, columns_to_analyze)

```

Value counts for scheme_management:

scheme_management	
VWC	36793
WUG	5206
Water authority	3153
WUA	2883
Water Board	2748
Parastatal	1680
Private operator	1063
Company	1061
Other	766
SWC	97
Trust	72

Name: count, dtype: int64

Value counts for management:

management	
vwc	40507
wug	6515
water board	2933
wua	2535
private operator	1971
parastatal	1768
water authority	904
other	844
company	685
unknown	561
other - school	99
trust	78

Name: count, dtype: int64

Value counts for management_group:

management_group	
user-group	52490
commercial	3638
parastatal	1768
other	943
unknown	561

```
Name: count, dtype: int64
```

Considering the project's context, it's evident that `scheme_management` and `management` are closely associated, both concerning the management of wells. However, `management_group` seems unrelated to the other two columns and has fewer unique values.

Given the prevalence of missing values in `scheme_management` compared to `management`, it's more practical to utilize `management` for further analysis.

To explore the relationship between `management_group` and `management`, a groupby operation will be conducted on the pandas dataframe. This will help ascertain how different management groups correspond to various management types, providing valuable insights for the project.

```
# Create a duplicate DataFrame and drop scheme_management column
new_data = data.copy().drop(['scheme_management'], axis=1)
```

```
# Group by management_group column and management column
pd.DataFrame(new_data.groupby(['management_group',
                              'management']).size())
```

management_group	management	0
commercial	company	685
	private operator	1971
	trust	78
	water authority	904
other	other	844
	other - school	99
	parastatal	1768
unknown	unknown	561
user-group	vwc	40507
	water board	2933
	wua	2535
	wug	6515

Since the "management_group" column has fewer different types compared to the "management" column, we're dropping the "management" column to keep things simpler. This helps avoid making our analysis and predictions too complicated, making it easier to understand and work with the data. By doing this, we still keep the important information about how wells are managed, but without unnecessary details that could confuse things.

```
# Drop management column
new_data = new_data.drop(['management'], axis=1)
```

3.1.2 extraction_type | extraction_type_group | extraction_type_class

The columns `extraction_type`, `extraction_type_group`, and `extraction_type_class` all pertain to the type of extraction used. An assessment of their value counts will be conducted to determine the similarity of the data.

```
# Define columns to analyze
columns_to_analyze = ['extraction_type', 'extraction_type_group',
                      'extraction_type_class']
```

```
# Call display_value_counts function
display_value_counts(data, columns_to_analyze)
```

Value counts for `extraction_type`:

extraction_type	
gravity	26780
nira/tanira	8154
other	6430
submersible	4764
swn 80	3670
mono	2865
india mark ii	2400
afridev	1770
ksb	1415
other - rope pump	451
other - swn 81	229
windmill	117
india mark iii	98
ceso	90
other - play pump	85
walimi	48
climax	32
other - mkulima/shinyanga	2

Name: count, dtype: int64

Value counts for `extraction_type_group`:

extraction_type_group	
gravity	26780
nira/tanira	8154
other	6430
submersible	6179
swn 80	3670
mono	2865
india mark ii	2400
afridev	1770
rope pump	451
other handpump	364
other motorpump	122
wind-powered	117

```
india mark iii      98
Name: count, dtype: int64
```

Value counts for extraction_type_class:

```
extraction_type_class
gravity      26780
handpump    16456
other       6430
submersible 6179
motorpump   2987
rope pump    451
wind-powered 117
Name: count, dtype: int64
```

it's noticeable that the "extraction_type" and "extraction_type_group" columns share similarities, representing various extraction methods. However, the "extraction_type_class" column appears unrelated to the other two, containing fewer unique values. Within the project context, it's apparent that these three columns redundantly convey information about extraction methods. Hence, we'll conduct a groupby operation on all three columns to assess their relationships and decide on the most suitable approach for analysis

```
# Group by extraction_type, extraction_type_group, and
extraction_type_class
```

```
pd.DataFrame(new_data.groupby(['extraction_type_class',
'extraction_type_group', 'extraction_type']).size())
```

```
0
extraction_type_class extraction_type_group extraction_type
gravity                gravity            gravity
26780
handpump              afridev            afridev
1770
                    india mark ii        india mark ii
2400
                    india mark iii       india mark iii
98
                    nira/tanira          nira/tanira
8154
                    other handpump        other - mkulima/shinyanga
2
                                                other - play pump
85
                                                other - swan 81
229
                                                walimi
```

48		
	swn 80	swn 80
3670		
motorpump	mono	mono
2865		
	other motorpump	ceмо
90		
		climax
32		
other	other	other
6430		
rope pump	rope pump	other - rope pump
451		
submersible	submersible	ksb
1415		
		submersible
4764		
wind-powered	wind-powered	windmill
117		

Given the analysis conducted, it's evident that the "extraction_type_group" column provides a more detailed classification of extraction methods compared to the "extraction_type_class" column. Additionally, "extraction_type_group" contains fewer unique values than the "extraction_type" column. Therefore, within the project's context, opting to utilize the "extraction_type_group" column allows for a balance between addressing the curse of dimensionality and retaining relevant information. This decision is informed by the need to streamline the dataset while preserving essential details about extraction methods.

```
# Drop extraction_type_class and extraction_type columns
new_data = new_data.drop(['extraction_type_class', 'extraction_type'],
axis=1)
```

3.1.3 payment | payment_type

The columns `payment` and `payment_type` both relate to monetary transactions. An examination of their value counts will be conducted to assess the similarity of data.

```
# define columns to analyze
columns_to_analyze = ['payment', 'payment_type']

# Call display_value_counts function
display_value_counts(data, columns_to_analyze )
```

Value counts for payment:

payment	
never pay	25348
pay per bucket	8985
pay monthly	8300
unknown	8157

```
pay when scheme fails    3914
pay annually             3642
other                   1054
Name: count, dtype: int64
```

```
Value counts for payment_type:
payment_type
never pay      25348
per bucket     8985
monthly        8300
unknown        8157
on failure     3914
annually       3642
other          1054
Name: count, dtype: int64
```

Since the `payment` and `payment_type` columns exhibit a perfect relationship, we'll remove the `payment` column as it contains more verbose information compared to the `payment_type` column.

```
# Drop payment column
new_data = new_data.drop(['payment'], axis=1)
```

3.1.4 `water_quality` | `quality_group`

The `water_quality` and `quality_group` columns both pertain to the quality of water from the well. We will assess their value counts to determine the similarity of data.

```
# Define columns to analyze
columns_to_analyze = ['water_quality', 'quality_group']

# Call display_value_counts function
display_value_counts(data, columns_to_analyze)
```

```
Value counts for water_quality:
water_quality
soft          50818
salty         4856
unknown       1876
milky          804
coloured       490
salty abandoned 339
fluoride       200
fluoride abandoned 17
Name: count, dtype: int64
```

```
Value counts for quality_group:
quality_group
good          50818
salty         5195
unknown       1876
milky         804
colored       490
fluoride      217
Name: count, dtype: int64
```

Given that the `water_quality` column has slightly more unique values compared to the `quality_group` column, though the disparity is not significant, we will drop the `quality_group` column due to its lower informational content.

```
# Drop the quantity_group column
new_data = new_data.drop(['quantity_group'], axis=1)
```

3.1.6 source | source_type | source_class

The columns `source`, `source_type`, and `source_class` all pertain to the source of water from the well. We will assess their value counts to determine the similarity of data.

```
# Define columns to analyze
columns_to_analyze = ['source', 'source_type', 'source_class']

# Call display_value_counts function
display_value_counts(data, columns_to_analyze)
```

```
Value counts for source:
source
spring          17021
shallow well    16824
machine dbh     11075
river           9612
rainwater harvesting 2295
hand dtw        874
lake            765
dam             656
other           212
unknown         66
Name: count, dtype: int64
```

```
Value counts for source_type:
source_type
spring          17021
```



```

shallow well          16824
borehole              11949
river/lake            10377
rainwater harvesting   2295
dam                   656
other                  278
Name: count, dtype: int64

```

```

Value counts for source_class:
source_class
groundwater    45794
surface        13328
unknown         278
Name: count, dtype: int64

```

Upon reviewing the value counts, we observe that the `source` column exhibits a greater diversity of unique values compared to the `source_type` column, and likewise, the `source_type` column contains more unique values than the `source_class` column. In the context of the project, we opt for the `source_type` column as it strikes a balance between reducing dimensionality and retaining relevant information. This choice is driven by the fact that the `source_type` column offers more detailed information than the `source_class` column while also presenting fewer unique values than the `source` column, thus aiding in simplifying the dataset without significant loss of information.

```

# Drop 'source' and 'source_class' columns
new_data = new_data.drop(['source', 'source_class'], axis=1)

```

3.1.7 waterpoint_type | waterpoint_type_group

The `waterpoint_type` and `waterpoint_type_group` are both related columns that talk about the type of the waterpoint. We shall begin by evaluating their value counts in order to check the similarity of data.

```

# Define the columns to analyze
columns_to_analyze = ['waterpoint_type', 'waterpoint_type_group']

# Call the display_value_counts function
display_value_counts(data, columns_to_analyze)

```

```

Value counts for waterpoint_type:
waterpoint_type
communal standpipe    28522
hand pump             17488
other                  6380
communal standpipe multiple  6103
improved spring       784

```

```

cattle trough      116
dam                 7
Name: count, dtype: int64

Value counts for waterpoint_type_group:
waterpoint_type_group
communal standpipe    34625
hand pump             17488
other                 6380
improved spring       784
cattle trough         116
dam                   7
Name: count, dtype: int64

```

Upon examining the value counts, it's evident that the `waterpoint_type` column encompasses a greater variety of unique values compared to the `waterpoint_type_group` column. In the context of the project, we opt to retain the `waterpoint_type` column due to its richer granularity of information compared to the `waterpoint_type_group` column. This decision is guided by the need to preserve detailed insights about waterpoint types while maintaining a manageable dataset.

```

# Drop waterpoint_type_group column
new_data = new_data.drop(['waterpoint_type_group'], axis=1)

```

3.1.8 Dropping Unnecessary Columns |

```

def drop_columns_and_print_reason(data, columns):
    """
    Drop specified columns from a DataFrame and print a message
    indicating the reason for dropping.

    Parameters:
    - data: DataFrame
        The DataFrame containing the data.
    - columns: list
        A list of column names to be dropped.
    """
    data.drop(columns=columns, axis=1, inplace=True)
    for column in columns:
        print(f"Dropped column '{column}' because it did not add any
value to our analysis.")
columns_to_drop = ['longitude', 'latitude', 'wpt_name', 'num_private',
                    'region_code', 'district_code', 'public_meeting',
                    'recorded_by', 'scheme_name']

```

```
drop_columns_and_print_reason(new_data, columns_to_drop)
```

Dropped column 'longitude' because it did not add any value to our analysis.

Dropped column 'latitude' because it did not add any value to our analysis.

Dropped column 'wpt_name' because it did not add any value to our analysis.

Dropped column 'num_private' because it did not add any value to our analysis.

Dropped column 'region_code' because it did not add any value to our analysis.

Dropped column 'district_code' because it did not add any value to our analysis.

Dropped column 'public_meeting' because it did not add any value to our analysis.

Dropped column 'recorded_by' because it did not add any value to our analysis.

Dropped column 'scheme_name' because it did not add any value to our analysis.

We have now finalized the selection of columns for our analysis. We will proceed to clean the data using the following columns:

Numerical Columns:

- `gps_height`
- `population`
- `amount_tsh`
- `date_recorded`
- `construction_year`

Categorical Columns:

- `funder`
- `installer`
- `basin`
- `subvillage`
- `region`
- `lga`
- `ward`
- `scheme_management`
- `permit`
- `extraction_type_group`
- `payment_type`
- `management_group`
- `water_quality`

- quantity
- source_type
- waterpoint_type
- status_group


3.2 Data Cleaning

In this phase, we will address missing values and identify duplicate records within the dataset.

3.2.1 visualizing the new Data set

In this section, we will be looking at the missing values in the dataset.

```
msno.bar(new_data, figsize=(15, 5), color='deepskyblue', fontsize=12)
# Changing color to ocean blue
plt.title("Missing Data Visualisation")
plt.xticks(rotation=90)
plt.show()
```



Based on the visualization above, the following columns contain missing data:

- funder
- installer
- subvillage
- permit

```
print(new_data.columns)
Index(['amount_tsh', 'date_recorded', 'funder', 'gps_height',
      'installer',
      'basin', 'subvillage', 'region', 'lga', 'ward', 'population',
      'permit',
      'construction_year', 'extraction_type_group',
```

```

'management_group',
  'payment_type', 'water_quality', 'quality_group', 'quantity',
  'source_type', 'waterpoint_type', 'status_group'],
  dtype='object')

def missing_values_summary(data, columns=None):
    """
    Calculate the total missing values and their percentages for
    specified columns in the DataFrame.

    Parameters:
    - data: DataFrame
        The DataFrame containing the data.
    - columns: list or None, optional (default=None)
        A list of columns to include in the summary. If None, all
        columns will be included.

    Returns:
    - summary: DataFrame
        A DataFrame summarizing the total missing values and their
        percentages for specified columns.
    """
    if columns is None:
        columns = data.columns
    total_missing = data[columns].isnull().sum()
    percent_missing = (total_missing / len(data)) * 100
    summary = pd.DataFrame({
        'Total Missing': total_missing,
        'Percentage Missing': percent_missing
    })
    summary = summary[summary['Total Missing'] > 0] # Exclude columns
    with no missing values
    return summary

missing_values_summary(new_data, columns=['funder', 'installer',
    'subvillage', 'permit'])

```

	Total Missing	Percentage Missing
funder	3637	6.122896
installer	3655	6.153199
subvillage	371	0.624579
permit	3056	5.144781

3.2.1.1 Dropping Missing Values

```

def drop_columns_below_threshold(data, columns, threshold=7):
    """
    Drop rows from DataFrame if their missing percentage is below the
    specified threshold.

    Parameters:

```

```

- data: DataFrame
    DataFrame containing the data.
- columns: list
    A list of column names to be evaluated for dropping.
- threshold: float, optional (default=7)
    threshold below which rows will be dropped.


Returns:
- data: DataFrame
    DataFrame with rows dropped if their missing percentage is
    below the threshold.
"""
for column in columns:
    missing_percentage = (data[column].isnull().sum() / len(data))
* 100
    if missing_percentage < threshold:
        data = data[data[column].notna()]
return data

# List of columns to evaluate
columns_to_evaluate = ['funder', 'installer', 'subvillage', 'permit']

# Drop columns if their missing percentage is below the threshold
new_data = drop_columns_below_threshold(new_data, columns_to_evaluate)

# Visualize the missing values in the new dataset
msno.bar(new_data, figsize=(15, 5), color='deepskyblue', fontsize=12)
# Changing color to ocean blue
plt.title("Missing Data Visualisation")
plt.xticks(rotation=90)
plt.show()

```



3.2.2 Check for Duplicates and Outliers

```
# Check for duplicated records
```

```
new_data[new_data.duplicated()]
```

	amount_tsh	date_recorded	funder	gps_height	\
id					
59310	0.0	2011-07-18	Government Of Tanzania	0	
2296	0.0	2011-07-31	Kkkt_makwale	0	
53399	0.0	2011-07-23	Danida	0	
2859	0.0	2011-07-31	Kkkt_makwale	0	
70330	0.0	2011-04-12	Rc Church	0	
...	
16666	0.0	2011-07-26	Md	0	
5064	0.0	2011-07-12	Hesawa	0	
47527	0.0	2011-04-12	Rc Church	0	
58255	0.0	2011-07-13	Do	0	
40607	0.0	2011-04-15	Government Of Tanzania	0	

	installer	basin	subvillage	region	
lga \					
id					
59310	Government	Lake Victoria	Nyanza	Mwanza	
Geita					
2296	KKKT _ Konde and DWE	Lake Nyasa	Isimba	Mbeya	
Kyela					
53399	Central government	Lake Nyasa	Malungo	Mbeya	
Kyela					
2859	KKKT _ Konde and DWE	Lake Nyasa	Isimba	Mbeya	
Kyela					
70330	RC Church	Lake Nyasa	Matwalani	Mbeya	Mbeya
Rural					
...	
...					
16666	DW	Lake Victoria	Kishoju 1	Kagera	
Muleba					
5064	HESAWA	Lake Victoria	Mishenye	Kagera	Bukoba
Rural					
47527	RC Church	Lake Nyasa	Mjimwema	Mbeya	Mbeya
Rural					
58255	DO	Lake Victoria	Kaluyango	Kagera	
Muleba					
40607	Government	Lake Rukwa	Mbuyuni A	Mbeya	
Chunya					

	ward	population	permit	construction_year	\
id					
59310	Kalangalala	0	True	0	
2296	Makwale	0	True	0	
53399	Mwaya	0	True	0	

2859	Makwale	0	True	0
70330	Ulenje	0	False	0
...
16666	Nshamba	0	True	0
5064	Buterankuzi	0	True	0
47527	Ulenje	0	False	0
58255	Ikondo	0	True	0
40607	Mbuyuni	0	True	0

extraction_type_group management_group payment_type				
water_quality \				
id				
59310	submersible	user-group	never pay	
soft				
2296	gravity	user-group	never pay	
soft				
53399	gravity	user-group	never pay	
soft				
2859	gravity	user-group	never pay	
soft				
70330	gravity	user-group	on failure	
soft				
...
.				
16666	submersible	user-group	never pay	
soft				
5064	gravity	user-group	monthly	
soft				
47527	gravity	user-group	on failure	
soft				
58255	gravity	user-group	never pay	
soft				
40607	gravity	user-group	never pay	
soft				

quality_group quantity source_type waterpoint_type \					
id					
59310	good	insufficient	borehole	communal	standpipe
2296	good	enough	spring	communal	standpipe
53399	good	dry	spring	communal	standpipe
2859	good	enough	spring	communal	standpipe
70330	good	enough	river/lake	communal	standpipe
...
16666	good	enough	spring	communal	standpipe
5064	good	enough	river/lake	communal	standpipe
47527	good	enough	river/lake	communal	standpipe
58255	good	enough	spring	communal	standpipe
40607	good	enough	spring	communal	standpipe

id		status_group
59310		functional
2296		functional
53399	non	functional
2859		functional
70330	non	functional
...		...
16666	non	functional
5064		functional
47527		functional
58255		functional
40607	non	functional

[2617 rows x 22 columns]

The presence of duplicated records in our dataset doesn't imply inaccuracies in the data. Rather, it suggests that multiple wells were constructed as part of the same project, resulting in identical features across these records.





3.2.2.1 Outliers

We'll address outliers in the numerical columns relevant to our project, including:

- amount_tsh
- gps_height
- population
- construction_year

```
# Function to visualize outliers as a box plot.
def visualize_outliers(df, col, title):
    """
    This function visualizes the outliers and outputs boxplots
    """
    plt.figure(figsize=(10, 2))
    sns.boxplot(x=col, data=df)
    plt.title(title, fontsize=15)
    plt.show()

visualize_outliers(data, 'amount_tsh', 'Amount TSH Outliers')
visualize_outliers(data, 'gps_height', 'GPS Height Outliers')
visualize_outliers(data, 'population', 'Population Outliers')
visualize_outliers(data, 'construction_year', 'Construction Year Outliers')
```



In the boxplot of the total static head column, we observe outliers, but these values do not necessarily indicate erroneous data, as a total static head of even 250,000 is possible. Similarly, although the majority of data is clustered around 0, further investigation is needed to understand this distribution in univariate analysis.

Regarding the `gps_height` column, no outliers are present, indicating no need for outlier treatment.

In the population column, outliers exist, yet they do not signify erroneous data, as populations of up to 30,000 near a well are plausible. This variable's values are region-dependent.

Concerning the construction year column, no outliers are evident. However, some nonsensical years may indicate data errors, warranting further investigation in univariate analysis. Overall, no outlier treatment is necessary for the construction year column.

3.2.3 Uniformity

In this section, we will be looking at the uniformity of the data. Uniformity refers to the consistency of the data with respect to the formatting, labelling. We will be looking at the following:

- Labelling
- Formatting

```
# function to rename columns
def rename_columns(df, col):
    """
    This function re-formats the column names to ensure that they are
    all in lower case and contain no spaces
    """
    new_col = col.strip().replace(' ', '_').lower()

    df.rename(columns = {col: new_col}, inplace = True)

    # Loop through all column names and reformat
    for col in new_data.columns:
        rename_columns(new_data, col)
```

```
new_data.columns
```

```
Index(['amount_tsh', 'date_recorded', 'funder', 'gps_height',
       'installer',
       'basin', 'subvillage', 'region', 'lga', 'ward', 'population',
       'permit',
       'construction_year', 'extraction_type_group',
       'management_group',
       'payment_type', 'water_quality', 'quality_group', 'quantity',
       'source_type', 'waterpoint_type', 'status_group'],
      dtype='object')
```

```
# Rename 'lga' column to 'local_government_area'
new_data.rename(columns = {'lga': 'local_government_area'}, inplace =
True)
```

```
new_data.columns
```

```
Index(['amount_tsh', 'date_recorded', 'funder', 'gps_height',
       'installer',
```

```

        'basin', 'subvillage', 'region', 'local_government_area',
'ward',
        'population', 'permit', 'construction_year',
'extraction_type_group',
        'management_group', 'payment_type', 'water_quality',
'quality_group',
        'quantity', 'source_type', 'waterpoint_type', 'status_group'],
dtype='object')

```

3.2.3.2 Formatting

```

# function to to convert the columns to the appropriate data types
def convert_to_category(df, col):
    """
    This function converts the object categories to category data type
    """
    df[col] = df[col].astype('category')

# for loop to convert the object categoris to a category datatype
for col in new_data.select_dtypes(include='object').columns:
    convert_to_category(new_data, col)

import numpy as np

# Replace "0" values in the 'construction_year' column with NaN
new_data['construction_year'] =
new_data['construction_year'].replace(0, np.nan)

# Convert the 'construction_year' column to datetime format
new_data['construction_year'] =
pd.to_datetime(new_data['construction_year'], format='%Y')

```

3.2.4 Data Consistency Check

This section focuses on assessing the consistency of the data.

We'll examine the consistency of the following columns:

- installer
- subvillage

```

def check_installer_consistency(data):
    """
    Check the consistency of the 'installer' column in the dataset.

    Parameters:
    data (DataFrame): The dataset containing the 'installer' column.

    Returns:
    list: A list of unique values in the 'installer' column.
    """

```

```

"""
unique_installers = data['installer'].unique()
return unique_installers

print(new_data.columns)

Index(['amount_tsh', 'date_recorded', 'funder', 'gps_height',
'installer',
      'basin', 'subvillage', 'region', 'local_government_area',
'ward',
      'population', 'permit', 'construction_year',
'extraction_type_group',
      'management_group', 'payment_type', 'water_quality',
'quality_group',
      'quantity', 'source_type', 'waterpoint_type', 'status_group'],
      dtype='object')

# Call the function to check consistency in the 'installer' column
installer_consistency = check_installer_consistency(new_data)
print(installer_consistency)

['Roman', 'GRUMETI', 'World vision', 'UNICEF', 'Artisan', ...,
'DWE/Anglican church', 'VIFAI', 'Dina', 'brown', 'SELEPTA']
Length: 2040
Categories (2040, object): ['- ', '0', 'A.D.B', 'AAR', ..., 'wizara ya
maji', 'world', 'world banks', 'world vision']

# From the most common 100 value counts we realized some spelling
mistakes or different syntax between same categories

# Replacing the spelling mistakes and collect same categories in same
name

new_data['installer'].replace(to_replace = ('District Water
Department', 'District water depar', 'Distric Water Department'),
                             value = 'District water department' ,
inplace=True)

new_data['installer'].replace(to_replace = ('FinW', 'Fini water', 'FINI
WATER'), value = 'Fini Water' , inplace=True)
new_data['installer'].replace(to_replace = 'JAICA', value = 'Jaica' ,
inplace=True)

new_data['installer'].replace(to_replace = ('COUN', 'District
COUNCIL', 'DISTRICT COUNCIL', 'District Counci',
      'District
Council', 'Council', 'Counc', 'District Council', 'Distri'),
                             value = 'District council' ,
inplace=True)
new_data
new_data['installer'].replace(to_replace = ('RC CHURCH', 'RC Churc',

```

```

'RC','RC Ch','RC C', 'RC CH','RC church',
                                'RC CATHORIC',) , value ='RC
Church' , inplace=True)

new_data['installer'].replace(to_replace = ('Central
Government','Tanzania Government',
                                'central government','Cental
Government', 'Cebtral Government',
                                'Tanzanian Government','Tanzania
government', 'Centra Government' ,
                                'CENTRAL GOVERNMENT', 'TANZANIAN
GOVERNMENT','Central govt', 'Centr',
                                'Centra govt')) , value ='Central
government' , inplace=True)

new_data['installer'].replace(to_replace = ('World vision', 'World
Division','World Vision'),
                                value ='world vision' ,
inplace=True)

new_data['installer'].replace(to_replace = ('Unisef','UNICEF'),value
='Unicef' , inplace=True)
new_data['installer'].replace(to_replace = 'DANID', value ='DANIDA' ,
inplace=True)

new_data['installer'].replace(to_replace = ('villigers', 'villager',
'Villagers', 'Villa', 'Village', 'Villi',
                                'Village Council','Village
Council', 'Villages', 'Vill', 'Village community',
                                'Villaers', 'Village Community',
'Villag','Village Council', 'Village council',
                                'Village Council','Villagerd',
'Villager', 'Village Technician',
                                'Village Office','Village
community members'),
                                value ='villagers' ,
inplace=True)

new_data['installer'].replace(to_replace
=('Commu','Communit','commu','COMMU', 'COMMUNITY') ,
                                value ='Community' ,
inplace=True)
new_data['installer'].replace(to_replace = ('GOVERNMENT', 'GOVER',
'GOVERNME', 'GOVERN', 'GOVERN', 'Gover', 'Gove',
                                'Governme','Governmen' ) ,value
='Government' , inplace=True)

new_data['installer'].replace(to_replace = 'Hesawa' ,value ='HESAWA' ,
inplace=True)

```

```

# continue to replacing spellin mistakes and getting together values
new_data['installer'].replace(to_replace = ('Colonial Government') ,
value = 'Colonial government' , inplace=True)
new_data['installer'].replace(to_replace = ('Government of Misri') ,
value = 'Misri Government' , inplace=True)
new_data['installer'].replace(to_replace = ('Italy government') ,
value = 'Italian government' , inplace=True)
new_data['installer'].replace(to_replace = ('British colonial
government') , value = 'British government' , inplace=True)
new_data['installer'].replace(to_replace = ('Concern /government') ,
value = 'Concern/Government' , inplace=True)
new_data['installer'].replace(to_replace = ('Village Government') ,
value = 'Village government' , inplace=True)
new_data['installer'].replace(to_replace = ('Government and
Community') , value = 'Government /Community' , inplace=True)
new_data['installer'].replace(to_replace = ('Cetral government /RC') ,
value = 'RC church/Central Gover' , inplace=True)
new_data['installer'].replace(to_replace = ('Government
/TCRS','Government/TCRS') , value = 'TCRS /Government' , inplace=True)
new_data['installer'].replace(to_replace = ('ADRA /Government') ,
value = 'ADRA/Government' , inplace=True)

# Sub village

# Drop data that is less than 3 characters and more than 20 characters
new_data.drop(new_data[new_data['subvillage'].str.len() < 3].index,
axis=0, inplace=True)

new_data[new_data['subvillage'].str.len() < 3]

Empty DataFrame
Columns: [amount_tsh, date_recorded, funder, gps_height, installer,
basin, subvillage, region, local_government_area, ward, population,
permit, construction_year, extraction_type_group, management_group,
payment_type, water_quality, quality_group, quantity, source_type,
waterpoint_type, status_group]
Index: []

new_data.to_csv('wells_cleaned_data.csv', index=False)

```

3.3 Exploratory Data Analysis

This section involves exploring the columns in the dataset. We will conduct:

- Univariate Analysis
- Bivariate Analysis
- Multivariate Analysis

3.4.1 Univariate Analysis

This section focuses on exploring individual columns in the dataset to analyze their distributions and gather insights. We'll divide this analysis into two main parts:

- Categorical Columns Examination
- Numerical Columns Examination

```
categorical_columns = new_data.select_dtypes(include='object').columns
categorical_columns
```

```
Index([], dtype='object')
```

```
numerical_columns = new_data.select_dtypes(include='number').columns
numerical_columns
```

```
Index(['amount_tsh', 'gps_height', 'population'], dtype='object')
```

Function to Both plot and get value counts of a column

```
import matplotlib.pyplot as plt
```

```
def get_value_counts(df, col):
```

```
    """
    Returns the value counts of a column in a dataframe including NaN
    values, sorted in descending order.
```

```
    Parameters:
```

```
        df (DataFrame): The DataFrame containing the data.
```

```
        col (str): The name of the column for which value counts are
    calculated.
```

```
    Returns:
```

```
        counts (Series): The value counts of the specified column.
```

```
    """
```

```
    counts = df[col].value_counts(dropna=False)
```

```
    return counts
```

```
def plot_data(df, col, title):
```

```
    """
    Plots the top 20 value counts of a column in a dataframe as a bar
    chart and provides a description of the plot.
```

```
    Parameters:
```

```
        df (DataFrame): The DataFrame containing the data.
```

```
        col (str): The name of the column to be visualized.
```

```
        title (str): The title of the plot.
```

```
    """
```

```
    top_20_counts = get_value_counts(df, col).head(20)
```

```
    plt.figure(figsize=(10, 5))
```

```
    top_20_counts.plot(kind='bar', color='#037bfc', edgecolor='black',
    fontsize=10)
```




```

plt.title(title)
plt.xlabel(col)
plt.ylabel('Count')
plt.xticks(rotation=90)
plt.show()

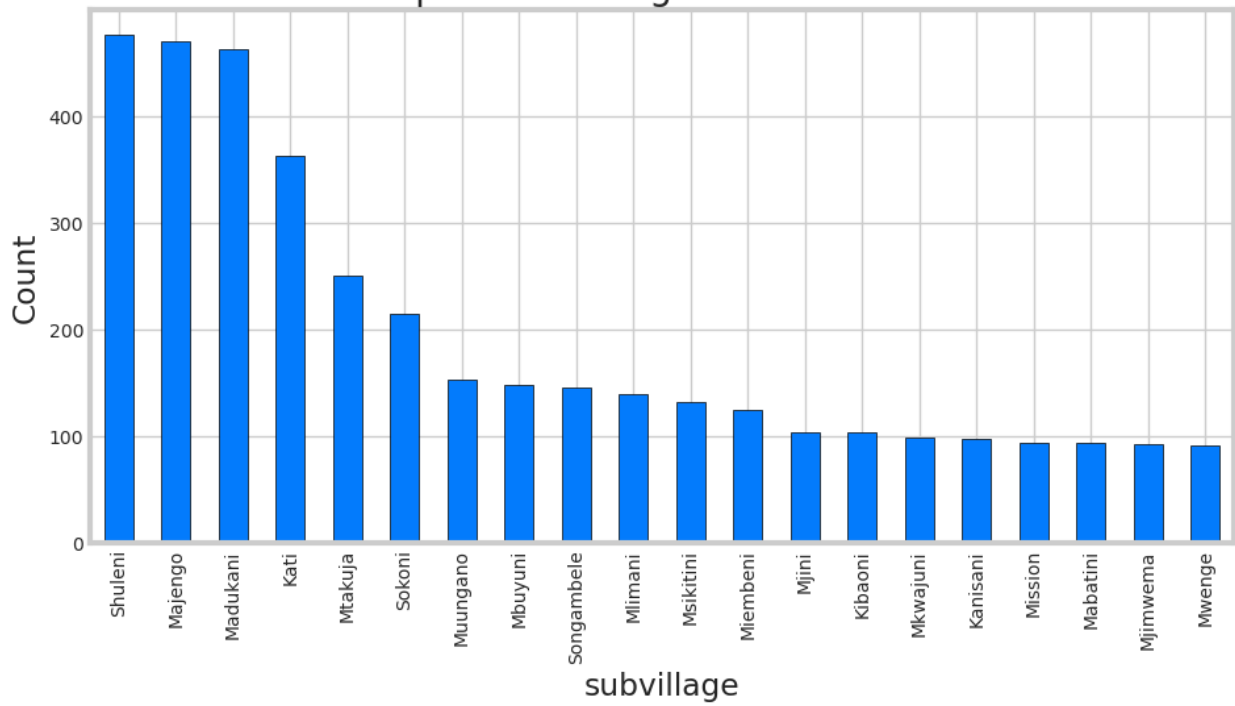
columns_to_plot = [ 'ward', 'subvillage',
                    'basin',]

for col in columns_to_plot:
    plot_data(new_data, col, f'Top 20 {col} Value Counts')

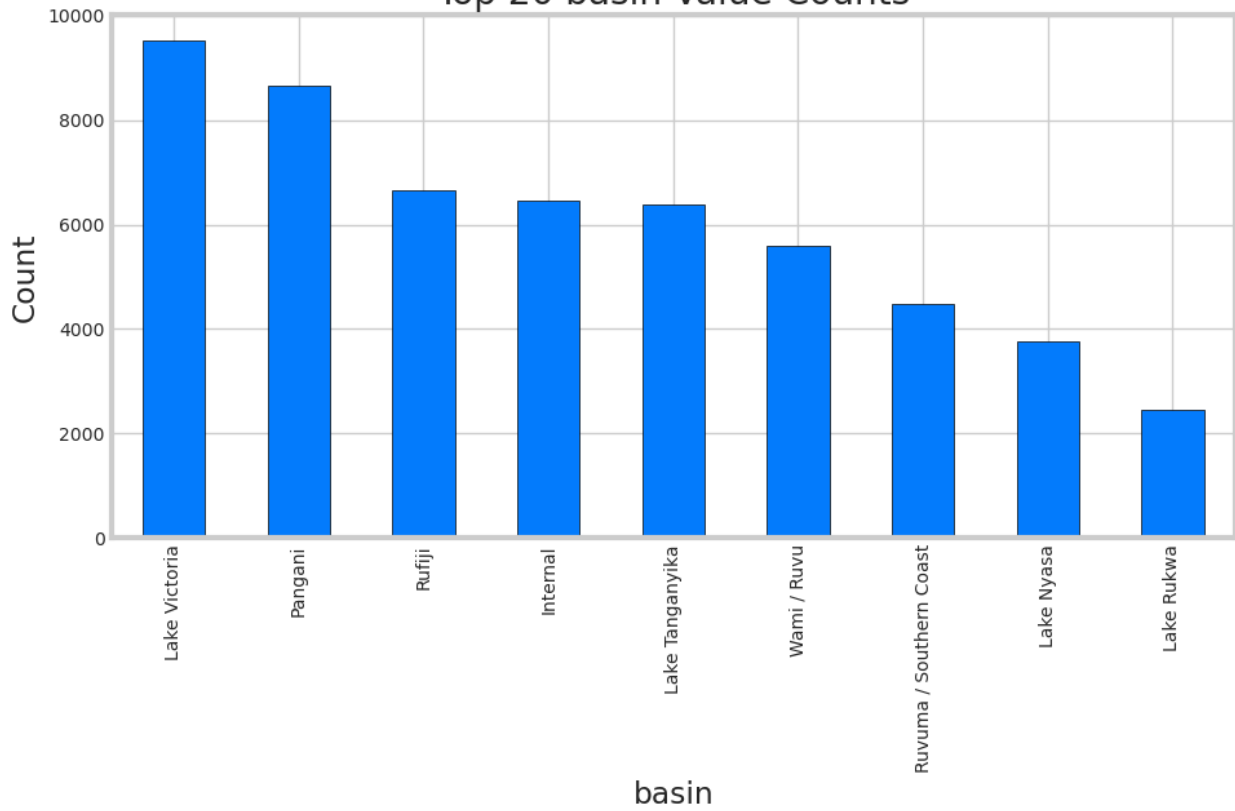
```




Top 20 subvillage Value Counts



Top 20 basin Value Counts




```
# Histogram for construction_year
plt.figure(figsize=(10, 6))
sns.histplot(data=new_data, x='construction_year', bins=30, kde=True)
plt.title('Distribution of Construction Years')
plt.xlabel('Construction Year')
plt.ylabel('Frequency')
plt.show()
```



3.4.1.2 Analysis of Numerical Columns

In this subsection, we will focus on the analysis of the dataset's numerical columns. There are a total of four numerical columns, and we will conduct univariate analysis on each of them individually.

```
## Comparing basin and functionality of wells columns
plt.figure(figsize=(18,10))
blue_palette = sns.color_palette("Blues", n_colors=3) # Using 'Blues'
palette with 3 different shades of blue
ax = sns.countplot(x='basin', hue="status_group", data=new_data,
palette=blue_palette)
plt.show()
```




```
# Define a colorful blue palette
blue_palette = sns.color_palette(['#4A90E2', '#6CC0E5', '#8ED8E2',
                                  '#AACFD0', '#D1D1D1'])

# Get the top 10 regions by count
top_regions = new_data['region'].value_counts().head(10).index


# Filter the data to include only the top 10 regions
filtered_data = new_data[new_data['region'].isin(top_regions)]

# Plot the countplot with the specified palette
plt.figure(figsize=(18, 10))
ax = sns.countplot(x='region', hue='status_group', data=filtered_data,
                  palette=blue_palette)
plt.show()
```



```
# Define a custom colorful blue palette
blue_palette = ['#4A90E2', '#6CC0E5', '#8ED8E2', '#AACFD0', '#D1D1D1',
                '#FFC0CB', '#87CEEB', '#00BFFF', '#1E90FF', '#4682B4']

# Plot the countplot with the specified palette
plt.figure(figsize=(18, 10))
ax = sns.countplot(x='basin', hue='quantity', data=new_data,
                  palette=blue_palette)
plt.show()
```



```
def corrmatrix(df):
    ''' This function plots a correlation matrix for numeric features
    in a given dataframe '''
    plt.figure(figsize=(10, 5))


    # Filter numeric columns
    numeric_cols = df.select_dtypes(include='number')

    # Calculate correlation matrix
    corr = numeric_cols.corr()

    # Generate a mask to only show the bottom triangle
    mask = np.triu(np.ones_like(corr, dtype=bool))

    # Generate heatmap
    sns.heatmap(round(corr, 2), annot=True, mask=mask, vmin=-1,
    vmax=1, cmap='Blues')
    plt.title('Correlation of Numeric Features')
    plt.show()

# Call the function
corrmatrix(new_data)
```



4. Modeling

The objective of this notebook is to construct a predictive model that can classify the status of Tanzanian wells based on the available dataset. We will evaluate the model

primarily based on its accuracy, aiming for a minimum threshold of 75% accuracy. To achieve this goal, we will employ several machine learning algorithms:

- Logistic Regression
- Decision Tree
- Random Forest
- XGBoost

Our evaluation will focus on accuracy as the primary metric. Additionally, to ensure the models generalize well and do not overfit, we will utilize cross-validation techniques during the evaluation process.

```
#view the new data
import pandas as pd
df = pd.read_csv('wells_cleaned_data.csv')
df.head()
```

	amount_tsh	date_recorded	funder	gps_height	installer \
0	6000.0	2011-03-14	Roman	1390	Roman
1	0.0	2013-03-06	Grumeti	1399	GRUMETI
2	25.0	2013-02-25	Lottery Club	686	world vision
3	0.0	2013-01-28	Unicef	263	Unicef
4	0.0	2011-07-13	Action In A	0	Artisan

	basin	subvillage	region	local_government_area
0	Lake Nyasa	Mnyusi B	Iringa	Ludewa
1	Lake Victoria	Nyamara	Mara	Serengeti
2	Pangani	Majengo	Manyara	Simanjiro
3	Ruvuma / Southern Coast	Mahakamani	Mtwara	Nanyumbu
4	Lake Victoria	Kyanyamisa	Kagera	Karagwe

	ward	population	permit	construction_year
0	Mundindi	109	False	1999-01-01
1	Natta	280	True	2010-01-01
2	Ngorika	250	True	2009-01-01
3	Nanyumbu	58	True	1986-01-01
4	Nyakasimbi	0	True	NaN

management_group	payment_type	water_quality	quality_group
------------------	--------------	---------------	---------------

quantity \	user-group	annually	soft	good
0	user-group	annually	soft	good
enough				
1	user-group	never pay	soft	good
insufficient				
2	user-group	per bucket	soft	good
enough				
3	user-group	never pay	soft	good
dry				
4	other	never pay	soft	good
seasonal				

	source_type	waterpoint_type	status_group
0	spring	communal standpipe	functional
1	rainwater harvesting	communal standpipe	functional
2	dam	communal standpipe multiple	functional
3	borehole	communal standpipe multiple	non functional
4	rainwater harvesting	communal standpipe	functional

```

#get copy of our dataframe
dfcopy = df.copy()

#check value count of target group
dfcopy['status_group'].value_counts(normalize=True)

status_group
functional          0.541287
non functional      0.387962
functional needs repair 0.070751
Name: proportion, dtype: float64

```

The "Functional" category constitutes the largest portion of the dataset, representing 54.1%. Meanwhile, the "Non-functional" category comprises 38.8% of the dataset, and the "Functional needs repair" category makes up 7% of the dataset.

```

new_status_group = {'non functional': 0, 'functional': 1, 'functional
needs repair' : 2}
dfcopy['status_group'] =
dfcopy['status_group'].replace(new_status_group)

dfcopy['status_group'].value_counts()

status_group
1    29210
0    20936
2     3818
Name: count, dtype: int64

```

The current data type is "object," necessitating conversion to an integer.


```
dfcopy['status_group'].dtypes
```

```
dtype('int64')
```

```
categorical = ['source_type', 'quantity', 'water_quality',  
'payment_type', 'management_group', 'basin']
```

```
ohe = pd.get_dummies(df[categorical], prefix=categorical,  
drop_first=True)  
ohe
```

	source_type_dam	source_type_other	source_type_rainwater
harvesting \			
0	False	False	
False			
1	False	False	
True			
2	True	False	
False			
3	False	False	
False			
4	False	False	
True			
...	
...			
53959	False	False	
False			
53960	False	False	
False			
53961	False	False	
False			
53962	False	False	
False			
53963	False	False	
False			

	source_type_river/lake	source_type_shallow well
source_type_spring \		
0	False	False
True		
1	False	False
False		
2	False	False
False		
3	False	False
False		
4	False	False
False		
...
...		

53959	False	False
False		
53960	False	False
True		
53961	True	False
False		
53962	False	True
False		
53963	False	True
False		

	quantity_enough	quantity_insufficient	quantity_seasonal	\
0	True	False	False	
1	False	True	False	
2	True	False	False	
3	False	False	False	
4	False	False	True	
...	
53959	True	False	False	
53960	True	False	False	
53961	True	False	False	
53962	False	True	False	
53963	True	False	False	

	quantity_unknown	water_quality_fluoride	\
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	False	
...	
53959	False	False	
53960	False	False	
53961	False	False	
53962	False	False	
53963	False	False	

	water_quality_fluoride	abandoned	water_quality_milky	\
0		False	False	
1		False	False	
2		False	False	
3		False	False	
4		False	False	
...		
53959		False	False	
53960		False	False	
53961		False	False	
53962		False	False	
53963		False	False	

	water_quality_salty	water_quality_salty	abandoned
water_quality_soft \			
0	False		False
True			
1	False		False
True			
2	False		False
True			
3	False		False
True			
4	False		False
True			
...
...			
53959	False		False
True			
53960	False		False
True			
53961	False		False
True			
53962	False		False
True			
53963	True		False
False			
	water_quality_unknown	payment_type_monthly	payment_type_never
pay \			
0	False	False	
False			
1	False	False	
True			
2	False	False	
False			
3	False	False	
True			
4	False	False	
True			
...
...			
53959	False	True	
False			
53960	False	False	
False			
53961	False	False	
False			
53962	False	False	
True			
53963	False	False	
False			

bucket	payment_type_on failure	payment_type_other	payment_type_per
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	False	
...	
53959	False	False	
53960	False	False	
53961	False	False	
53962	False	False	
53963	True	False	

	payment_type_unknown	management_group_other	\
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	True	
...	
53959	False	False	
53960	False	False	
53961	False	False	
53962	False	False	
53963	False	False	

	management_group_parastatal	management_group_unknown	\
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	False	
...	
53959	False	False	
53960	False	False	
53961	False	False	
53962	False	False	

53963	False	False
management_group_user-group	basin_Lake Nyasa	basin_Lake Rukwa
\		
0	True	False
1	True	False
2	True	False
3	True	False
4	False	False
...
53959	True	False
53960	True	False
53961	True	False
53962	True	False
53963	True	False
basin_Lake Tanganyika	basin_Lake Victoria	basin_Pangani
0	False	False
1	False	True
2	False	False
3	False	True
4	False	False
...
53959	False	False
53960	False	True
53961	False	False
53962	False	False
53963	False	False
basin_Rufiji	basin_Ruvuma / Southern Coast	basin_Wami / Ruvu
0	False	False
1	False	False
2	False	False
3	False	True
4	False	False

...
53959	False	False	True
53960	False	False	False
53961	True	False	False
53962	True	False	False
53963	False	False	True

[53964 rows x 35 columns]

Preview the one hot encoded datatypes
 ohe.dtypes

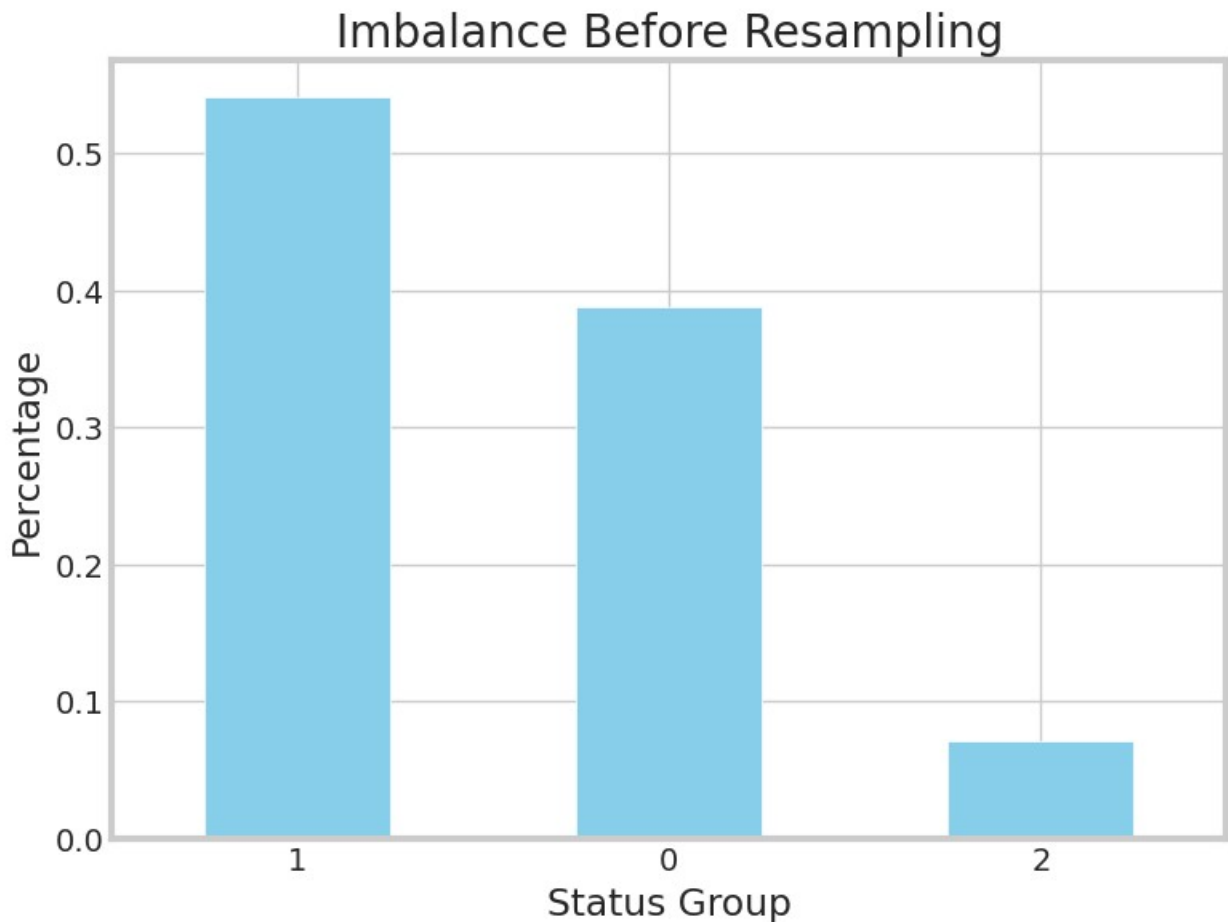
source_type_dam	bool
source_type_other	bool
source_type_rainwater harvesting	bool
source_type_river/lake	bool
source_type_shallow well	bool
source_type_spring	bool
quantity_enough	bool
quantity_insufficient	bool
quantity_seasonal	bool
quantity_unknown	bool
water_quality_fluoride	bool
water_quality_fluoride abandoned	bool
water_quality_milky	bool
water_quality_salty	bool
water_quality_salty abandoned	bool
water_quality_soft	bool
water_quality_unknown	bool
payment_type_monthly	bool
payment_type_never pay	bool
payment_type_on failure	bool
payment_type_other	bool
payment_type_per bucket	bool
payment_type_unknown	bool
management_group_other	bool
management_group_parastatal	bool
management_group_unknown	bool
management_group_user-group	bool
basin_Lake Nyasa	bool
basin_Lake Rukwa	bool
basin_Lake Tanganyika	bool
basin_Lake Victoria	bool

```
basin_Pangani          bool
basin_Rufiji           bool
basin_Ruvuma / Southern Coast  bool
basin_Wami / Ruvu      bool
dtype: object
```

Check for Class Imbalance

```
import matplotlib.pyplot as plt

# Check imbalance before resampling
plt.figure(figsize=(8, 6))
dfcopy['status_group'].value_counts(normalize=True).plot(kind='bar',
color='skyblue')
plt.title('Imbalance Before Resampling')
plt.xlabel('Status Group')
plt.ylabel('Percentage')
plt.xticks(rotation=0)
plt.show()
```



```

# Define X and y
X = ohe # Assuming you have already performed one-hot encoding
y = dfcopy['status_group']

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=1)

# Instantiate logistic regression
logreg = LogisticRegression(class_weight='balanced', solver='lbfgs',
random_state=42)

# Build a pipeline with standard scaler and logistic regression
scaled_pipeline = Pipeline([('ss', StandardScaler()),
                             ('logreg', logreg)])

# Fit the training data to pipeline
scaled_pipeline.fit(X_train, y_train)

# Predict the labels of the test set: y_pred
y_pred_log = scaled_pipeline.predict(X_test)

# Perform cross-validation: cv_results
cv_results_log = cross_validate(scaled_pipeline, X_test, y_test, cv=3)

# Display cross-validation test scores
print("Cross Validation Test Scores:", cv_results_log['test_score'])


# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred_log)

# Plot confusion matrix
plt.figure(figsize=(10, 10))
target_names = ['Non-Functional', 'Functional', 'Needs-Repair']
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=target_names, yticklabels=target_names)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

# Classification report
target_names = ['Non-Functional', 'Functional', 'Needs-Repair']
print(classification_report(y_test, y_pred_log,
target_names=target_names))

Cross Validation Test Scores: [0.56003335 0.55336298 0.54879066]

```

	precision	recall	f1-score	support
Non-Functional	0.70	0.56	0.62	4175
Functional	0.75	0.55	0.64	5851
Needs-Repair	0.15	0.60	0.24	767
accuracy			0.56	10793
macro avg	0.53	0.57	0.50	10793
weighted avg	0.69	0.56	0.60	10793

4.2 Decision Tree

Decision trees are a type of non-parametric supervised learning algorithm utilized for both classification and regression tasks. Their objective is to construct a model capable of predicting the target variable's value by discerning straightforward decision rules derived from the features within the data.

```
# Define X and y
X = ohe # Assuming you have already performed one-hot encoding
y = dfcopy['status_group']

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=1)

# Instantiate the Decision Tree Classifier
clf = DecisionTreeClassifier(criterion='entropy', random_state=39)

# Build a pipeline with standard scaler and decision tree classifier
tree_pipeline = Pipeline([
    ('ss', StandardScaler()), # Standard scaling (not really needed
for decision trees, but included for consistency)
    ('clf', clf) # Decision Tree Classifier
])

# Fit train data to pipeline
tree_pipeline.fit(X_train, y_train)

# Predict the labels of the test set: y_pred
y_pred_dt = tree_pipeline.predict(X_test)

# Perform cross-validation: cv_results
cv_results_dt = cross_validate(tree_pipeline, X_test, y_test, cv=3)

# Display cross-validation test scores
print("Cross Validation Test Scores:", cv_results_dt['test_score'])


# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred_dt)

# Plot confusion matrix
plt.figure(figsize=(10, 10))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=target_names, yticklabels=target_names)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Decision Tree Confusion Matrix')
plt.show()

# Classification report
target_names = ['Non-Functional', 'Functional', 'Needs-Repair']
```

```
print(classification_report(y_test, y_pred_dt,
target_names=target_names))
```

Cross Validation Test Scores: [0.67898833 0.67954419 0.67528496]



	precision	recall	f1-score	support
Non-Functional	0.71	0.61	0.66	4175
Functional	0.70	0.83	0.76	5851
Needs-Repair	0.43	0.14	0.21	767
accuracy			0.70	10793

macro avg	0.62	0.53	0.54	10793
weighted avg	0.69	0.70	0.68	10793

4.3 Random Forest

Random forests, also known as random decision forests, are a powerful ensemble learning technique used for classification, regression, and other tasks. They work by creating a large number of decision trees during training and then combining their predictions to determine the final outcome. For classification tasks, the mode of the classes predicted by individual trees is chosen, while for regression tasks, the mean prediction of the individual trees is taken.

```
# Define X and y
X = ohe # Assuming you have already performed one-hot encoding
y = dfcopy['status_group']

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=1)

# Build a pipeline with StandardScaler and RandomForestClassifier
random_pipeline = Pipeline([('ss', StandardScaler()),
('RF',
RandomForestClassifier(random_state=0))])

# Fit the training data to pipeline
random_pipeline.fit(X_train, y_train)

# Predict the labels of the test set: y_pred_rf
y_pred_rf = random_pipeline.predict(X_test)

# Perform cross-validation: cv_results_rf
cv_results_rf = cross_validate(random_pipeline, X_test, y_test, cv=3)


# Display cross-validation test scores
print("Cross Validation Test Scores:", cv_results_rf['test_score'])

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred_rf)

# Plot confusion matrix
plt.figure(figsize=(10, 10))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=target_names, yticklabels=target_names)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Random Forest Confusion Matrix')
plt.show()
```

```
# Classification report
target_names = ['Non-Functional', 'Functional', 'Needs-Repair']
print(classification_report(y_test, y_pred_rf,
target_names=target_names))
```

Cross Validation Test Scores: [0.68343524 0.68037799 0.67973311]



	precision	recall	f1-score	support
Non-Functional	0.72	0.60	0.66	4175
Functional	0.70	0.84	0.76	5851
Needs-Repair	0.44	0.14	0.22	767

accuracy			0.70	10793
macro avg	0.62	0.53	0.55	10793
weighted avg	0.69	0.70	0.68	10793

4.5 XGBOOST

XGBoost is a highly efficient and flexible gradient boosting library that is optimized for distributed computing. It is designed to be portable and offers implementation of machine learning algorithms within the Gradient Boosting framework. XGBoost employs parallel tree boosting techniques, also referred to as GBDT or GBM, which enables solving various data science problems rapidly and accurately.

```
# Instantiate XGBClassifier
clf = XGBClassifier()

# Define X and y
X = ohe # Assuming you have already performed one-hot encoding
y = dfcopy['status_group']

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=1)

# Build a pipeline with StandardScaler and XGBClassifier
xgb_pipeline = Pipeline([
    ('ss', StandardScaler()),
    ('xgb', clf)
])

# Fit XGBClassifier
xgb_pipeline.fit(X_train, y_train)

# Predict on training and test sets
training_preds = xgb_pipeline.predict(X_train)
test_preds = xgb_pipeline.predict(X_test)

# Accuracy of training and test sets
training_accuracy = accuracy_score(y_train, training_preds)
test_accuracy = accuracy_score(y_test, test_preds)

print('Training Accuracy: {:.4}%'.format(training_accuracy * 100))
print('Validation accuracy: {:.4}%'.format(test_accuracy * 100))

# Define parameter grid for GridSearchCV
param_grid = {
    'xgb__learning_rate': [0.1, 0.2],
    'xgb__max_depth': [6],
    'xgb__min_child_weight': [1, 2],
```

```

    'xgb__subsample': [0.5, 0.7],
    'xgb__n_estimators': [100],
}

# Perform GridSearchCV
grid_clf = GridSearchCV(xgb_pipeline, param_grid, scoring='accuracy',
cv=None, n_jobs=1)
grid_clf.fit(X_train, y_train)

# Get best parameters
best_parameters = grid_clf.best_params_

print('Grid Search found the following optimal parameters: ')
for param_name in sorted(best_parameters.keys()):
    print('%s: %r' % (param_name, best_parameters[param_name]))

# Predict on training and test sets using best model
training_preds = grid_clf.predict(X_train)
test_preds = grid_clf.predict(X_test)

# Accuracy of training and test sets using best model
training_accuracy = accuracy_score(y_train, training_preds)
test_accuracy = accuracy_score(y_test, test_preds)

print('')
print('Training Accuracy: {:.4}%'.format(training_accuracy * 100))
print('Validation accuracy: {:.4}%'.format(test_accuracy * 100))

Training Accuracy: 71.78%
Validation accuracy: 70.17%
Grid Search found the following optimal parameters:
xgb__learning_rate: 0.2
xgb__max_depth: 6
xgb__min_child_weight: 1
xgb__n_estimators: 100
xgb__subsample: 0.7

Training Accuracy: 71.61%
Validation accuracy: 70.06%

```

Evaluation

In this project, we employed several machine learning algorithms to predict the status of Tanzanian wells. Here are the results:

Logistic Regression

- The logistic regression model achieved an accuracy of 56%. While this performance served as a baseline, it fell short of our desired accuracy of 75%.

- The confusion matrix revealed a bias towards predicting that a well is functional. However, the model demonstrated balanced performance with a higher number of true positives and true negatives than false positives and false negatives.
- Cross-validation consistently returned stable scores, indicating that the model wasn't overfitting the data.

Decision Tree

- The decision tree model improved upon the baseline, reaching an accuracy of 70%.
- Similar to logistic regression, the confusion matrix showed a bias towards predicting functional wells. The model displayed balanced performance, indicating no overfitting.
- Cross-validation results remained consistent across tests.

Random Forest

- The random forest model also achieved an accuracy of 70%, matching the performance of the decision tree model.
- Confusion matrix analysis and cross-validation yielded similar observations as with the decision tree model, indicating balanced performance.
- XGBoost
- After training the XGBoost model, we achieved a training accuracy of 71.78% and a validation accuracy of 70.17%.
- Grid search was employed to find the optimal hyperparameters, resulting in the following settings:
 - Learning Rate: 0.2
 - Maximum Depth: 6
 - Minimum Child Weight: 1
 - Number of Estimators: 100
 - Subsample: 0.7
- While the training and validation accuracies are slightly lower than those of other models, XGBoost still demonstrates robust performance and generalization capabilities.

How Models Were Evaluated

- Used a pipeline to scale the data and fit it to each model, followed by cross-validation to assess generalization performance.
- Confusion matrices provided insights into the models' prediction biases and overall performance.
- While models fell short of the desired 75% accuracy, achieving 70% accuracy indicates a promising start for the project.

Limitations

Despite our achievements, we encountered several limitations:

- Approximately 30% of the data remained unaccounted for by our models.

- Class imbalance across the output variables (functional, non-functional, needs-repair) may have affected model performance.

Conclusion

In this project, we applied several machine learning algorithms to predict the operational status of Tanzanian wells. Despite achieving an accuracy of around 70%, which is a decent baseline, there is room for improvement. Logistic regression, decision trees, and random forests were among the algorithms we explored, each offering unique insights into the data.

While our models performed reasonably well, there are several limitations and areas for improvement. The class imbalance across the three output variables (functional, non-functional, needs-repair) posed a challenge, potentially affecting the models' predictive accuracy.

Next Steps

Moving forward, there are several steps we can take to enhance the predictive capabilities of our models:

- **Investigate Additional Features:** Concentrating on geographical indicators like climate, population, and amount of water available in the area.
- **Time-Series Analysis:** Further consideration of the well ages should be analyzed to predict the average lifetime of more robust well structures.
- **Repairs:** Local governments should look at what type of water wells are needing repairs, and the severity of those repairs, to fine-tune non-functional indicators.